

```
*****
* HFSC Scheduler — Hierarchical + RT heap + Link-sharing (paper-faithful)
* DPDK dataplane implementation
*
* Implements:
* - Leaf deadlines (SCED-style) with eligible times (Eq. 7, 11)
* - Link-sharing via hierarchical virtual times (Eq. 12), SSF policy
* - Two min-heaps:
*   * eligible_heap: among eligible leaves, pick min deadline
*   * pending_heap: non-eligible leaves keyed by eligible_time
* - Piecewise-linear service curves (m1, m2, x) per paper §IV
*
* Notes:
* - Deadlines/eligible times are in wall-clock cycles (TSC)
* - Virtual times are in “service-normalized” cycles (via inverse curves)
* - Leaf classes guarantee service curves; interior classes use virtual times
* - Simplified piecewise-linear inverse computations for efficiency
******/
```

```
#include <stdint.h>
#include <stdbool.h>
#include <stdlib.h>
#include <string.h>

#include <rte_mbuf.h>
#include <rte_ring.h>
```

```
#include <rte_cycles.h>
#include <rte_ip.h>
#include <rte_tcp.h>
#include <rte_udp.h>
#include <rte_ether.h>
#include <rte_byteorder.h>

/* ===== CONFIG ===== */

#define QUEUE_SIZE 8192
#define MAX_CHILDREN 8
#define MAX_CLASSES 64
#define HEAP_CAPACITY MAX_CLASSES

/* ===== SERVICE CURVE ===== */
/* Two-piece linear curve:
 * - First segment slope m1 (bytes/sec) up to x bytes
 * - Second segment slope m2 (bytes/sec) after x
 */
typedef struct {
    uint64_t m1; /* bytes/sec */
    uint64_t m2; /* bytes/sec */
    uint64_t x; /* bytes */
} service_curve_t;

/* ===== HFSC CLASS ===== */
```

```

typedef struct hfsc_class {
    struct hfsc_class *parent;
    struct hfsc_class *children[MAX_CHILDREN];
    int num_children;

    bool is_leaf;
    struct rte_ring *q; /* only for leaf */

    /* Service curve */
    service_curve_t sc;

    /* Leaf RT state (wall-clock cycles) */
    uint64_t eligible_time; /* e_i */
    uint64_t deadline; /* d_i */

    /* Leaf RT accounting (bytes) */
    uint64_t rt_service; /* c_i: bytes served under RT criterion */

    /* Virtual time state (service-normalized cycles) */
    uint64_t vtime; /* v_i */
    uint64_t total_service; /* w_i: bytes served total (RT + LS) */

    /* Virtual curve parameters (piecewise-linear inverse):
     * Vi(a_k; v) = min(prev, S_i(v - v_parent + w_i(a_k)))
     * We store slopes and offset to compute inverse quickly.
    */
}

```

```
*/  
  
uint64_t vc_m1; /* slope in bytes/sec for first segment (maps v->bytes) */  
  
uint64_t vc_m2; /* slope in bytes/sec for second segment */  
  
uint64_t vc_x; /* breakpoint in “virtual domain” (cycles) */  
  
uint64_t vc_off; /* offset (bytes) accumulated at activation */
```

```
bool active;  
}  
hfsc_class_t;
```

```
/* ===== HEAP (min-heap) ===== */
```

```
typedef struct {  
  
    uint64_t key;  
  
    hfsc_class_t *cl;  
}  
heap_entry_t;
```

```
typedef struct {  
  
    heap_entry_t *entries;  
  
    int size;  
  
    int capacity;  
}  
min_heap_t;
```

```
static void heap_init(min_heap_t *h, int cap) {  
  
    h->entries = (heap_entry_t *)calloc(cap, sizeof(heap_entry_t));  
  
    h->size = 0;  
  
    h->capacity = cap;
```

```
}
```

```
static void heap_push(min_heap_t *h, uint64_t key, hfsc_class_t *cl) {
    if (h->size >= h->capacity) return;

    int i = h->size++;
    h->entries[i].key = key;
    h->entries[i].cl = cl;

    while (i > 0) {
        int p = (i - 1) / 2;
        if (h->entries[p].key <= h->entries[i].key) break;

        heap_entry_t tmp = h->entries[p];
        h->entries[p] = h->entries[i];
        h->entries[i] = tmp;

        i = p;
    }
}
```

```
static bool heap_empty(min_heap_t *h) { return h->size == 0; }
```

```
static heap_entry_t heap_top(min_heap_t *h) { return h->entries[0]; }
```

```
static heap_entry_t heap_pop(min_heap_t *h) {
    heap_entry_t min = h->entries[0];
    h->entries[0] = h->entries[--h->size];
    int i = 0;
    while (1) {
```

```

int l = 2 * i + 1, r = 2 * i + 2, s = i;

if (l < h->size && h->entries[l].key < h->entries[s].key) s = l;
if (r < h->size && h->entries[r].key < h->entries[s].key) s = r;
if (s == i) break;

heap_entry_t tmp = h->entries[s];
h->entries[s] = h->entries[i];
h->entries[i] = tmp;

i = s;

}

return min;
}

```

```
/* ===== GLOBALS ===== */
```

```

static hfsc_class_t *root;

static hfsc_class_t *site1, *site2;

static hfsc_class_t *site1_udp, *site1_tcp, *site2_udp, *site2_tcp;

```

```

static min_heap_t eligible_heap; /* key = deadline */

static min_heap_t pending_heap; /* key = eligible_time */

```

```
/* ===== TIME HELPERS ===== */
```

```
static inline uint64_t now_cycles(void) { return rte_get_tsc_cycles(); }
```

```
static inline uint64_t bytes_to_cycles(uint64_t bytes, uint64_t rate_bytes_per_sec) {
```

```

if (!rate_bytes_per_sec) return 0;

return (bytes * rte_get_tsc_hz()) / rate_bytes_per_sec;

}

/* ===== SERVICE CURVE HELPERS ===== */

/* Deadline increment for a packet of length len under SCED:

* If c_i + len <= x: use m1; else use m2.

*/
static inline uint64_t sc_deadline_delta(const service_curve_t *sc, uint64_t c_bytes, uint32_t
len_bytes) {

    uint64_t before = c_bytes;

    uint64_t after = c_bytes + len_bytes;

    if (after <= sc->x) {

        return bytes_to_cycles(len_bytes, sc->m1);

    } else if (before >= sc->x) {

        return bytes_to_cycles(len_bytes, sc->m2);

    } else {

        /* crosses breakpoint: split */

        uint64_t first = sc->x - before;

        uint64_t second = after - sc->x;

        return bytes_to_cycles(first, sc->m1) + bytes_to_cycles(second, sc->m2);

    }

}

```

/\* Eligible curve  $E_i(a_k; t)$ : for concave ( $m1 \geq m2$ ) equals  $D_i$ ; for convex, cap by  $m2$ .  
\* We implement the efficient rule from §IV:

```

* - concave: eligible == deadline
* - convex: eligible grows at m2 slope (i.e., cap RT allocation to asymptotic rate)
*/
static inline bool sc_is_concave(const service_curve_t *sc) { return sc->m1 >= sc->m2; }

/* ===== CLASSIFICATION ===== */
static inline hfsc_class_t *hfsc_classify(struct rte_mbuf *m) {
    struct rte_ipv4_hdr *ip = rte_pktmbuf_mtod_offset(m, struct rte_ipv4_hdr *, sizeof(struct rte_ether_hdr));
    if (!ip || ip->version != 4) return NULL;

    uint32_t src = rte_be_to_cpu_32(ip->src_addr);
    uint32_t dst = rte_be_to_cpu_32(ip->dst_addr);
    /* 192.168.2.20 -> 192.168.2.30 */
    if (src != 0xCOA80214 || dst != 0xCOA8021E) return NULL;

    if (ip->next_proto_id == IPPROTO_UDP) {
        struct rte_udp_hdr *udp = (struct rte_udp_hdr *)(ip + 1);
        uint16_t dport = rte_be_to_cpu_16(udp->dst_port);
        if (dport == 5001) return site1_udp;
        if (dport == 6001) return site2_udp;
    } else if (ip->next_proto_id == IPPROTO_TCP) {
        struct rte_tcp_hdr *tcp = (struct rte_tcp_hdr *)(ip + 1);
        uint16_t dport = rte_be_to_cpu_16(tcp->dst_port);
        if (dport == 5002) return site1_tcp;
    }
}

```

```

    if (dport == 6002) return site2_tcp;
}

return NULL;
}

/* ===== VIRTUAL CURVE UPDATE (Eq. 12) ===== */
/* We maintain virtual curve inverse parameters:
 *  $V_i(a_k; v) = \min(\text{prev}, S_i(v - v_{\text{parent}}(a_k)) + w_i(a_k))$ 
 * For piecewise-linear  $S_i$ , we set:
 * -  $vc\_m1 = sc.m1$ ,  $vc\_m2 = sc.m2$ 
 * -  $vc\_x = v_{\text{break}}$  = cycles corresponding to x bytes at slope m1
 * -  $vc\_off = w_i(a_k)$  (bytes already served)
 *  $v_{\text{parent}}$  is folded into vtime initialization at activation.
*/
static void update_virtual_curve(hfsc_class_t *cl, uint64_t v_parent, uint64_t w_bytes) {
    cl->vc_m1 = cl->sc.m1;
    cl->vc_m2 = cl->sc.m2;
    cl->vc_x = bytes_to_cycles(cl->sc.x, cl->sc.m1);
    cl->vc_off = w_bytes;
    /* Initialize  $v_i$  to  $\max(v_i, v_{\text{parent}})$  to keep siblings bounded (SSF with  $vs=(min+max)/2$  approx) */
    if (cl->vtime < v_parent) cl->vtime = v_parent;
}

/* Advance virtual time by inverse of  $S_i$  for bytes 'len':
 * If  $v < vc\_x$ : use m1; else use m2.

```

```

*/
static inline uint64_t vc_advance_cycles(hfsc_class_t *cl, uint32_t len_bytes) {
    /* Map bytes to virtual cycles using current segment slope.
     * We approximate by using m2 for fairness stability (bounded discrepancy),
     * but if cl->vtime < vc_x, we use m1 until crossing.
    */
    uint64_t v = cl->vtime;
    if (v < cl->vc_x) {
        /* first segment */
        uint64_t delta = bytes_to_cycles(len_bytes, cl->vc_m1);
        uint64_t newv = v + delta;
        if (newv <= cl->vc_x) return delta;
        /* crosses breakpoint: split */
        uint64_t first = cl->vc_x - v;
        uint64_t first_bytes = (first * cl->vc_m1) / rte_get_tsc_hz();
        uint64_t second_bytes = len_bytes - first_bytes;
        return first + bytes_to_cycles(second_bytes, cl->vc_m2);
    } else {
        return bytes_to_cycles(len_bytes, cl->vc_m2);
    }
}

/* ===== ACTIVATION/PASSIVE ===== */

static void set_active(hfsc_class_t *cl, uint64_t now) {
    if (cl->active) return;

```

```

cl->active = true;

if (cl->is_leaf) {

    /* Initialize RT state for new active period */

    cl->rt_service = 0;

    /* First packet's eligible/deadline:
     * Eligible: concave -> equals deadline; convex -> cap by m2 slope.
     * We start at now; next packet's increments computed on dequeue.
     */

    cl->eligible_time = now; /* eligible immediately for concave; for convex we gate via
pending_heap */

    cl->deadline = now; /* will be advanced on dequeue by sc_deadline_delta */

    /* Insert into appropriate heap */

    if (sc_is_concave(&cl->sc)) {

        heap_push(&eligible_heap, cl->deadline, cl);

    } else {

        /* convex: not eligible until we allocate at m2 rate; start pending at now */

        heap_push(&pending_heap, cl->eligible_time, cl);

    }

}

/* Initialize virtual curve at activation */

uint64_t v_parent = cl->parent ? cl->parent->vtime : 0;

update_virtual_curve(cl, v_parent, cl->total_service);

```

```

/* Bubble up activation */
if (cl->parent) set_active(cl->parent, now);

}

static void set_passive(hfsc_class_t *cl) {
    if (!cl->active) return;
    cl->active = false;

/* Bubble up: if parent has no active children, make it passive */
hfsc_class_t *p = cl->parent;
while (p) {
    bool any = false;
    for (int i = 0; i < p->num_children; i++) {
        if (p->children[i]->active) { any = true; break; }
    }
    if (any) break;
    p->active = false;
    p = p->parent;
}
}

/* ===== ENQUEUE ===== */
int hfsc_packet_in(struct rte_mbuf *m) {
    hfsc_class_t *cl = hfsc_classify(m);

```

```
    if (!cl || !cl->is_leaf) { rte_pktmbuf_free(m); return -1; }

    if (rte_ring_enqueue(cl->q, m) < 0) { rte_pktmbuf_free(m); return -1; }

    if (!cl->active) set_active(cl, now_cycles());
    return 0;
}
```

```
/* ===== RT HEAP MAINTENANCE ===== */
```

```
static void promote_pending(uint64_t now) {
    while (!heap_empty(&pending_heap) && heap_top(&pending_heap).key <= now) {
        heap_entry_t e = heap_pop(&pending_heap);
        if (e.cl->active) heap_push(&eligible_heap, e.cl->deadline, e.cl);
    }
}
```

```
/* ===== LINK-SHARING SELECTION (SSF) ===== */
```

```
static hfsc_class_t *ls_select(hfsc_class_t *cl) {
    if (!cl || !cl->active) return NULL;
    if (cl->is_leaf) return cl;

    hfsc_class_t *best = NULL;
    uint64_t best_v = UINT64_MAX;
    for (int i = 0; i < cl->num_children; i++) {
```

```

    hfsc_class_t *c = cl->children[i];

    if (!c->active) continue;

    if (c->vtime < best_v) { best_v = c->vtime; best = c; }

}

return ls_select(best);

}

```

```
/* ===== SCHEDULER ===== */
```

```

static hfsc_class_t *hfsc_select(uint64_t now) {

promote_pending(now);

if (!heap_empty(&eligible_heap)) {

    heap_entry_t top = heap_top(&eligible_heap);

    if (top.cl->active && top.key <= now) return top.cl;

}

return ls_select(root);

}

```

```
/* ===== ACCOUNTING ===== */
```

```

static void account_packet(hfsc_class_t *leaf, struct rte_mbuf *m, uint64_t now, bool used_rt) {

    uint32_t len = rte_pktmbuf_pkt_len(m);

/* Update leaf totals */

    leaf->total_service += len;

```

```

/* Virtual time advance at leaf */

leaf->vtime += vc_advance_cycles(leaf, len);

/* Propagate virtual time and totals up the tree */

hfsc_class_t *p = leaf->parent;

while (p) {

    p->total_service += len;

    /* SSF: keep vtime monotone and bounded among siblings */

    uint64_t dv = vc_advance_cycles(p, len);

    if (p->vtime < leaf->vtime) p->vtime = leaf->vtime; /* tighten to child min */

    p->vtime += dv;

    p = p->parent;

}

/* RT deadline/eligible updates if RT was used */

if (used_rt) {

    uint64_t delta = sc_deadline_delta(&leaf->sc, leaf->rt_service, len);

    leaf->rt_service += len;

    leaf->deadline += delta;

    if (sc_is_concave(&leaf->sc)) {

        /* eligible grows with deadline for concave */

        leaf->eligible_time = leaf->deadline;

    } else {

        /* convex: eligible grows at m2 slope only */

    }

}

```

```

leaf->eligible_time += bytes_to_cycles(len, leaf->sc.m2);

}

}

}

/* ===== DEQUEUE ===== */

struct rte_mbuf *hfsc_packet_out(void) {

    uint64_t now = now_cycles();

    hfsc_class_t *cl = hfsc_select(now);

    if (!cl || !cl->is_leaf) return NULL;

    /* Pop from eligible heap if RT path */

    bool used_rt = false;

    if (!heap_empty(&eligible_heap) && heap_top(&eligible_heap).cl == cl &&
heap_top(&eligible_heap).key <= now) {

        (void)heap_pop(&eligible_heap);

        used_rt = true;
    }

    struct rte_mbuf *m;

    if (rte_ring_dequeue(cl->q, (void **)&m) < 0) return NULL;

    account_packet(cl, m, now, used_rt);

    /* If queue still has packets, update next head's deadline/eligible */

```

```

if (!rte_ring_empty(cl->q)) {

    /* Peek next packet length: DPDK ring has no peek; we approximate by using typical MTU
       or defer until dequeue.

        * For correctness, we conservatively schedule next deadline on next dequeue (lazy
       update).

    */

    /* Reinsert into appropriate heap */

    if (sc_is_concave(&cl->sc)) {

        heap_push(&eligible_heap, cl->deadline, cl);

    } else {

        heap_push(&pending_heap, cl->eligible_time, cl);

    }

} else {

    set_passive(cl);

}

return m;

}

/* ===== CLASS CREATION ===== */

static hfsc_class_t *mk_class(hfsc_class_t *parent, service_curve_t sc, bool leaf, const char
*qname) {

    hfsc_class_t *c = (hfsc_class_t *)calloc(1, sizeof(*c));

    c->parent = parent;

    c->sc = sc;

    c->is_leaf = leaf;
}

```

```

if (parent) parent->children[parent->num_children++] = c;

if (leaf) {
    c->q = rte_ring_create(qname, QUEUE_SIZE, rte_socket_id(), RING_F_SP_ENQ | RING_F_SC_DEQ);
}

return c;
}

/* ===== INIT ===== */

void hfsc_init(void) {
    heap_init(&eligible_heap, HEAP_CAPACITY);
    heap_init(&pending_heap, HEAP_CAPACITY);

    /* Root: 100 Mbps */
    root = mk_class(NULL, (service_curve_t){ .m1 = 12500000, .m2 = 12500000, .x = 0 }, false, NULL);

    /* site1/site2: 50 Mbps each */
    site1 = mk_class(root, (service_curve_t){ .m1 = 6250000, .m2 = 6250000, .x = 0 }, false, NULL);
    site2 = mk_class(root, (service_curve_t){ .m1 = 6250000, .m2 = 6250000, .x = 0 }, false, NULL);

    /* site1 leaves: UDP 25 Mbps, TCP 25 Mbps (example split) */
    site1_udp = mk_class(site1, (service_curve_t){ .m1 = 25000000, .m2 = 25000000, .x = 0 }, true, "s1_udp");
}

```

```

site1_tcp = mk_class(site1, (service_curve_t){ .m1 = 25000000, .m2 = 25000000, .x = 0 }, true,
"s1_tcp");

/* site2 leaves: UDP 25 Mbps, TCP 25 Mbps */

site2_udp = mk_class(site2, (service_curve_t){ .m1 = 25000000, .m2 = 25000000, .x = 0 },
true, "s2_udp");

site2_tcp = mk_class(site2, (service_curve_t){ .m1 = 25000000, .m2 = 25000000, .x = 0 }, true,
"s2_tcp");

/* Initialize virtual times to zero */

root->vtime = site1->vtime = site2->vtime = 0;

site1_udp->vtime = site1_tcp->vtime = site2_udp->vtime = site2_tcp->vtime = 0;

}

/* ===== OPTIONAL: EXPORTS ===== */

hfsc_class_t *hfsc_root(void) { return root; }

hfsc_class_t *hfsc_site1(void) { return site1; }

hfsc_class_t *hfsc_site2(void) { return site2; }

hfsc_class_t *hfsc_site1_udp(void) { return site1_udp; }

hfsc_class_t *hfsc_site1_tcp(void) { return site1_tcp; }

hfsc_class_t *hfsc_site2_udp(void) { return site2_udp; }

hfsc_class_t *hfsc_site2_tcp(void) { return site2_tcp; }

//*****|2fwd integration*****

```

```

hfsc_init();

while (!force_quit) {

    // RX burst

    struct rte_mbuf *pkts_burst[MAX_PKT_BURST];
    uint16_t nb_rx = rte_ether_rx_burst(portid, queueid, pkts_burst, MAX_PKT_BURST);

    for (int i = 0; i < nb_rx; i++) {
        hfsc_packet_in(pkts_burst[i]);
    }

    // Scheduler + TX

    struct rte_mbuf *m;
    int nb_tx = 0;
    struct rte_mbuf *tx_pkts[MAX_PKT_BURST];
    while ((m = hfsc_packet_out()) != NULL && nb_tx < MAX_PKT_BURST) {
        tx_pkts[nb_tx++] = m;
    }
    if (nb_tx > 0) {
        rte_ether_tx_burst(portid, queueid, tx_pkts, nb_tx);
    }
}

```