

```
#ifndef HFSC_H
#define HFSC_H

#include <rte_mbuf.h>
#include <rte_ring.h>

#define QUEUE_SIZE 8192
#define MAX_CHILDREN 4
#define AVG_PKT_LEN 1500

typedef struct {
    uint64_t m1; // bytes/sec - initial slope
    uint64_t d; // us - delay for first segment (0 for linear)
    uint64_t m2; // bytes/sec - asymptotic slope
} service_curve_t;

typedef struct {
    double x; // start time (sec)
    double y; // start bytes
    double sm1; // slope 1 (bytes/sec)
    double sm2; // slope 2
    double dx; // x-length of first segment
    double dy; // y-length of first segment
} runtime_sc_t;

typedef struct hfsc_class {
```

```
struct hfsc_class *parent;

struct hfsc_class *children[MAX_CHILDREN];

int num_children;

bool is_leaf;

struct rte_ring *q;          // leaf only

service_curve_t rsc;
service_curve_t fsc;
service_curve_t usc;

runtime_sc_t deadline;
runtime_sc_t eligible;
runtime_sc_t virtual;
runtime_sc_t ulimit;

uint64_t cumul;             // RT service (bytes)
uint64_t total;              // total service (bytes)

uint64_t cl_e;               // eligible time (cycles)
uint64_t cl_d;               // deadline (cycles)
uint64_t cl_vt;              // virtual time (cycles)
uint64_t cl_myf;             // my fit time (from USC)
uint64_t cl_f;               // final fit time = max(myf, cfmin)

uint32_t vtperiod;
```

```
    uint32_t parentperiod;

    bool active;

    uint64_t last_time;

} hfsc_class_t;

extern hfsc_class_t *root, *site1, *site2, *udp1, *tcp1, *udp2, *tcp2;

void hfsc_init(void);

int hfsc_packet_in(struct rte_mbuf *m);

struct rte_mbuf *hfsc_packet_out(void);

#endif /* HFSC_H */

//*****
#include "hfsc.h"
//*****



* DPDK HFSC – Enhanced, Logically Faithful to Linux Kernel HFSC
* Features: USC (upper limit), accurate vt period, better peek
*****/


#include <stdint.h>
#include <stdbool.h>
#include <stdlib.h>
#include <math.h>
#include <rte_mbuf.h>
#include <rte_ring.h>
```

```

#include <rte_cycles.h>
#include <rte_ip.h>
#include <rte_udp.h>
#include <rte_tcp.h>
#include <rte_ether.h>
#include <rte_byteorder.h>

/* ===== CONFIG ===== */
#define QUEUE_SIZE 8192
#define MAX_CHILDREN 4
#define AVG_PKT_LEN 1500 // fallback when peek fails

/* ===== SERVICE CURVE ===== */
typedef struct {
    uint64_t m1; // bytes/sec - initial slope
    uint64_t d; // us - delay for first segment (0 for linear)
    uint64_t m2; // bytes/sec - asymptotic slope
} service_curve_t;

/* ===== RUNTIME SERVICE CURVE ===== */
typedef struct {
    double x; // start time (sec)
    double y; // start bytes
    double sm1; // slope 1 (bytes/sec)
    double sm2; // slope 2
    double dx; // x-length of first segment
}

```

```
    double dy; // y-length of first segment  
}  
  
/* ===== HFSC CLASS ===== */  
  
typedef struct hfsc_class {  
    struct hfsc_class *parent;  
    struct hfsc_class *children[MAX_CHILDREN];  
    int num_children;  
    bool is_leaf;  
  
    struct rte_ring *q; // leaf only  
  
    service_curve_t rsc; // real-time (deadline)  
    service_curve_t fsc; // fair/link-sharing  
    service_curve_t usc; // upper limit (optional)  
  
    runtime_sc_t deadline; // runtime D  
    runtime_sc_t eligible; // runtime E  
    runtime_sc_t virtual; // runtime V  
    runtime_sc_t ulimit; // runtime U (for USC)  
  
    uint64_t cumul; // RT service (bytes)  
    uint64_t total; // total service (bytes)  
  
    uint64_t cl_e; // eligible time (cycles)  
    uint64_t cl_d; // deadline (cycles)
```

```

    uint64_t cl_vt;           // virtual time (cycles)
    uint64_t cl_myf;          // my fit time (from USC)
    uint64_t cl_f;            // final fit time = max(myf, cfmin)

    uint32_t vtperiod;        // current virtual time period
    uint32_t parentperiod;    // parent's period when activated

    bool active;
    uint64_t last_time;       // last update/activation time
} hfsc_class_t;

/* ===== GLOBAL STATE ===== */
static hfsc_class_t *root;
static hfsc_class_t *site1, *site2;
static hfsc_class_t *udp1, *tcp1, *udp2, *tcp2;

/* ===== HELPERS ===== */
static inline uint64_t now_cycles(void) {
    return rte_get_tsc_cycles();
}

static inline double cycles_to_sec(uint64_t c) {
    return (double)c / rte_get_tsc_hz();
}

static inline double bytes_per_sec_to_per_cycle(uint64_t bps) {

```

```

    return (double)bps / rte_get_tsc_hz();

}

/* ===== RUNTIME CURVE MATH ===== */
static double rtsc_x2y(runtime_sc_t *rt, double x) {
    if (x <= rt->x) return rt->y;
    if (x <= rt->x + rt->dx)
        return rt->y + (x - rt->x) * rt->sm1;
    return rt->y + rt->dy + (x - rt->x - rt->dx) * rt->sm2;
}

static double rtsc_y2x(runtime_sc_t *rt, double y) {
    if (y <= rt->y) return rt->x;
    if (y <= rt->y + rt->dy)
        return rt->x + (y - rt->y) / rt->sm1;
    return rt->x + rt->dx + (y - rt->y - rt->dy) / rt->sm2;
}

static void rtsc_min(runtime_sc_t *rt, double new_x, double new_y,
                     double sm1, double sm2, double dx) {
    double y1 = rtsc_x2y(rt, new_x);
    double dy_new = dx * sm1;

    if (sm1 <= sm2) { // convex
        if (y1 < new_y) return;
        rt->x = new_x; rt->y = new_y; rt->dx = dx; rt->dy = dy_new;
    }
}

```

```

rt->sm1 = sm1; rt->sm2 = sm2;

return;
}

// concave

double y2 = rtsc_x2y(rt, new_x + dx);

if (y2 <= new_y + dy_new) {

    rt->x = new_x; rt->y = new_y; rt->dx = dx; rt->dy = dy_new;

    rt->sm1 = sm1; rt->sm2 = sm2;

    return;
}

// intersect - approximate new dx

double diff = y1 - new_y;

double dsm = sm1 - sm2;

double new_dx = (dsm > 0) ? diff / dsm : dx;

double new_dy = new_dx * sm1;

rt->x = new_x; rt->y = new_y; rt->dx = new_dx; rt->dy = new_dy;

rt->sm1 = sm1; rt->sm2 = sm2;

}

/* ===== PEEK NEXT PACKET LENGTH ===== */

static uint32_t peek_next_len(struct rte_ring *ring) {

    if (rte_ring_empty(ring)) return AVG_PKT_LEN;
}

```

```

void *obj;

if (rte_ring_dequeue(ring, &obj) == 0) {
    uint32_t len = rte_pktmbuf_pkt_len((struct rte_mbuf *)obj);
    rte_ring_enqueue(ring, obj); // put back
    return len;
}

return AVG_PKT_LEN;
}

/* ===== ACTIVATE / INIT CURVES ===== */
static void init_runtime_curve(runtime_sc_t *rt, double now_sec, double start_bytes,
                               uint64_t m1, uint64_t m2, uint64_t d) {
    double sm1 = bytes_per_sec_to_per_cycle(m1);
    double sm2 = bytes_per_sec_to_per_cycle(m2);
    double dx = (double)d / 1000000.0; // us → sec

    rt->x = now_sec;
    rt->y = start_bytes;
    rt->sm1 = sm1;
    rt->sm2 = sm2;
    rt->dx = dx;
    rt->dy = dx * sm1;
}

static void hfsc_activate(hfsc_class_t *cl, uint64_t now) {
    if (cl->active) return;
}

```

```

cl->active = true;

cl->last_time = now;

double now_sec = cycles_to_sec(now);

// Real-time curve

if (cl->rsc.m1 > 0 || cl->rsc.m2 > 0) {

    init_runtime_curve(&cl->deadline, now_sec, cl->cumul,
                       cl->rsc.m1, cl->rsc.m2, cl->rsc.d);

    cl->eligible = cl->deadline;

}

// Convex → eligible becomes linear m2

if (cl->rsc.m1 <= cl->rsc.m2) {

    cl->eligible.dx = 0;

    cl->eligible.dy = 0;

    cl->eligible.sm1 = bytes_per_sec_to_per_cycle(cl->rsc.m2);

    cl->eligible.sm2 = cl->eligible.sm1;

}

uint32_t next_len = peek_next_len(cl->q);

cl->cl_e = (uint64_t)(rtsc_y2x(&cl->eligible, cl->cumul) * rte_get_tsc_hz());

cl->cl_d = (uint64_t)(rtsc_y2x(&cl->deadline, cl->cumul + next_len) * rte_get_tsc_hz());

}

// Link-sharing curve

if (cl->fsc.m1 > 0 || cl->fsc.m2 > 0) {

```

```

    init_runtime_curve(&cl->virtual, now_sec, cl->total,
                       cl->fsc.m1, cl->fsc.m2, cl->fsc.d);

    cl->cl_vt = (uint64_t)(rtsc_y2x(&cl->virtual, cl->total) * rte_get_tsc_hz());

}

// Upper limit curve (USC)

if (cl->usc.m1 > 0 || cl->usc.m2 > 0) {

    init_runtime_curve(&cl->ulimit, now_sec, cl->total,
                       cl->usc.m1, cl->usc.m2, cl->usc.d);

    cl->cl_myf = (uint64_t)(rtsc_y2x(&cl->ulimit, cl->total) * rte_get_tsc_hz());

} else {

    cl->cl_myf = UINT64_MAX; // no limit

}

// VT period handling

cl->vtperiod++;

if (cl->parent) {

    cl->parentperiod = cl->parent->vtperiod;

}

if (cl->parent) hfsc_activate(cl->parent, now);

}

/* ===== CLASSIFICATION ===== */

static inline hfsc_class_t *hfsc_classify(struct rte_mbuf *m) {

```

```

struct rte_ipv4_hdr *ip = rte_pktmbuf_mtod_offset(m, struct rte_ipv4_hdr *, sizeof(struct
rte_ether_hdr));

if (ip->version != 4) return NULL;

uint32_t src_ip = rte_be_to_cpu_32(ip->src_addr);
uint32_t dst_ip = rte_be_to_cpu_32(ip->dst_addr);

if (src_ip != 0xc0a80214 || dst_ip != 0xc0a8021e) return NULL; // 192.168.2.20 → .30

uint8_t proto = ip->next_proto_id;
uint16_t dst_port = 0;

if (proto == IPPROTO_UDP) {

    struct rte_udp_hdr *udp = (struct rte_udp_hdr *)(ip + 1);

    dst_port = rte_be_to_cpu_16(udp->dst_port);

    if (dst_port == 5001) return udp1;
    if (dst_port == 6001) return udp2;

} else if (proto == IPPROTO_TCP) {

    struct rte_tcp_hdr *tcp = (struct rte_tcp_hdr *)(ip + 1);

    dst_port = rte_be_to_cpu_16(tcp->dst_port);

    if (dst_port == 5002) return tcp1;
    if (dst_port == 6002) return tcp2;

}

return NULL;
}

```

```
/* ===== ENQUEUE ===== */
```

```

int hfsc_packet_in(struct rte_mbuf *m) {
    hfsc_class_t *cl = hfsc_classify(m);
    if (!cl || !cl->is_leaf) {
        rte_pktmbuf_free(m);
        return -1;
    }

    uint64_t now = now_cycles();
    if (!cl->active) hfsc_activate(cl, now);

    if (rte_ring_enqueue(cl->q, m) < 0) {
        rte_pktmbuf_free(m);
        return -1;
    }
}

return 0;
}

/* ===== RT SELECT ===== */
static hfsc_class_t *hfsc_rt_select(uint64_t now) {
    hfsc_class_t *candidates[] = {udp1, tcp1, udp2, tcp2};
    hfsc_class_t *best = NULL;
    uint64_t min_d = UINT64_MAX;

    for (int i = 0; i < 4; i++) {
        hfsc_class_t *c = candidates[i];

```

```

if (c->active && c->cl_e <= now && c->cl_d < min_d) {

    min_d = c->cl_d;

    best = c;

}

return best;

}

/* ===== LS SELECT ===== */

static hfsc_class_t *hfsc_ls_select(hfsc_class_t *cl, uint64_t now) {

    if (!cl->active) return NULL;

    if (cl->is_leaf) return cl;

    // Find child with min cl_f (fit time)

    hfsc_class_t *best = NULL;

    uint64_t min_f = UINT64_MAX;

    for (int i = 0; i < cl->num_children; i++) {

        hfsc_class_t *child = cl->children[i];

        if (child->active && child->cl_f < min_f) {

            min_f = child->cl_f;

            best = child;

        }

    }

    if (!best) return NULL;

```

```

    return hfsc_ls_select(best, now);

}

/* ===== DEQUEUE & ACCOUNTING ===== */
struct rte_mbuf *hfsc_packet_out(void) {
    uint64_t now = now_cycles();
    if (!root->active) return NULL;

    hfsc_class_t *cl = hfsc_rt_select(now);
    bool is_realtime = (cl != NULL);

    if (!cl)
        cl = hfsc_ls_select(root, now);

    if (!cl || !cl->is_leaf) return NULL;

    struct rte_mbuf *m;
    if (rte_ring_dequeue(cl->q, (void **)&m) < 0)
        return NULL;

    uint32_t len = rte_pktmbuf_pkt_len(m);

    cl->total += len;
    if (is_realtime)
        cl->cumul += len;
}

```

```

double now_sec = cycles_to_sec(now);

// Update virtual curve

rtsc_min(&cl->virtual, now_sec, cl->total,
         bytes_per_sec_to_per_cycle(cl->fsc.m1),
         bytes_per_sec_to_per_cycle(cl->fsc.m2), 0);

cl->cl_vt = (uint64_t)(rtsc_y2x(&cl->virtual, cl->total) * rte_get_tsc_hz());

// Update USC (upper limit)

if (cl->usc.m1 > 0 || cl->usc.m2 > 0) {

    rtsc_min(&cl->ulimit, now_sec, cl->total,
              bytes_per_sec_to_per_cycle(cl->usc.m1),
              bytes_per_sec_to_per_cycle(cl->usc.m2), 0);

    cl->cl_myf = (uint64_t)(rtsc_y2x(&cl->ulimit, cl->total) * rte_get_tsc_hz());

}

// Update RT if applicable

if (cl->rsc.m1 > 0 || cl->rsc.m2 > 0) {

    uint32_t next_len = peek_next_len(cl->q);
    rtsc_min(&cl->deadline, now_sec, cl->cumul,
              bytes_per_sec_to_per_cycle(cl->rsc.m1),
              bytes_per_sec_to_per_cycle(cl->rsc.m2),
              (double)cl->rsc.d / 1000000.0);

    cl->eligible = cl->deadline;

    if (cl->rsc.m1 <= cl->rsc.m2) {

```

```

    cl->eligible.dx = 0;

    cl->eligible.dy = 0;

}

cl->cl_e = (uint64_t)(rtsc_y2x(&cl->eligible, cl->cumul) * rte_get_tsc_hz());
cl->cl_d = (uint64_t)(rtsc_y2x(&cl->deadline, cl->cumul + next_len) * rte_get_tsc_hz());

}

if (rte_ring_empty(cl->q)) {

    cl->active = false;

    cl->vtperiod++; // new period when reactivated

}

return m;

}

/* ===== INIT ===== */

void hfsc_init(void) {

    // Root - 100 Mbps (no USC)

    root = calloc(1, sizeof(*root));

    root->rsc = (service_curve_t){12500000, 0, 12500000};

    root->fsc = root->rsc;

    root->usc = (service_curve_t){0, 0, 0}; // no limit

    // site1 - 50 Mbps (with USC cap at 60 Mbps)

    site1 = calloc(1, sizeof(*site1));
}

```

```

site1->parent = root;

site1->rsc = (service_curve_t){6250000, 0, 6250000};

site1->fsc = site1->rsc;

site1->usc = (service_curve_t){7500000, 0, 7500000};

root->children[root->num_children++] = site1;

// udp1 - concave RT + USC

udp1 = calloc(1, sizeof(*udp1));

udp1->parent = site1;

udp1->rsc = (service_curve_t){5000000, 10000, 1250000};

udp1->fsc = (service_curve_t){1250000, 0, 1250000};

udp1->usc = (service_curve_t){2000000, 0, 2000000};

udp1->is_leaf = true;

udp1->q = rte_ring_create("udp1_q", QUEUE_SIZE, rte_socket_id(), RING_F_SP_ENQ | RING_F_SC_DEQ);

site1->children[site1->num_children++] = udp1;

// tcp1 - linear

tcp1 = calloc(1, sizeof(*tcp1));

tcp1->parent = site1;

tcp1->rsc = (service_curve_t){5000000, 0, 5000000};

tcp1->fsc = tcp1->rsc;

tcp1->usc = (service_curve_t){6000000, 0, 6000000};

tcp1->is_leaf = true;

tcp1->q = rte_ring_create("tcp1_q", QUEUE_SIZE, rte_socket_id(), RING_F_SP_ENQ | RING_F_SC_DEQ);

site1->children[site1->num_children++] = tcp1;

```

```

// site2 - 50 Mbps (with USC cap at 60 Mbps)

site2 = calloc(1, sizeof(*site2));

site2->parent = root;

site2->rsc = (service_curve_t){6250000, 0, 6250000};

site2->fsc = site2->rsc;

site2->usc = (service_curve_t){7500000, 0, 7500000};

root->children[root->num_children++] = site2;

// udp2 - concave RT + USC

udp2 = calloc(1, sizeof(*udp2));

udp2->parent = site2;

udp2->rsc = (service_curve_t){5000000, 10000, 1250000};

udp2->fsc = (service_curve_t){1250000, 0, 1250000};

udp2->usc = (service_curve_t){2000000, 0, 2000000};

udp2->is_leaf = true;

udp2->q = rte_ring_create("udp2_q", QUEUE_SIZE, rte_socket_id(), RING_F_SP_ENQ | RING_F_SC_DEQ);

site2->children[site2->num_children++] = udp2;

// tcp2 - linear

tcp2 = calloc(1, sizeof(*tcp2));

tcp2->parent = site2;

tcp2->rsc = (service_curve_t){5000000, 0, 5000000};

tcp2->fsc = tcp2->rsc;

tcp2->usc = (service_curve_t){6000000, 0, 6000000};

```

```
tcp2->is_leaf = true;

tcp2->q = rte_ring_create("tcp2_q", QUEUE_SIZE, rte_socket_id(), RING_F_SP_ENQ | RING_F_SC_DEQ);

site2->children[site2->num_children++] = tcp2;

}
```

```
////////////////////////////////////////////////////////////////////////
```

```
#include <stdio.h>

#include <stdint.h>

#include <signal.h>

#include <rte_eal.h>

#include <rte_ethdev.h>

#include <rte_mbuf.h>

#include <rte_mempool.h>

#include <rte_cycles.h>

#include "hfsc.h"

#define RX_RING_SIZE    1024

#define TX_RING_SIZE    1024

#define NUM_MBUFS      8191

#define MBUF_CACHE_SIZE 250

#define BURST_SIZE      32

static volatile bool force_quit = false;

static void signal_handler(int signum) {
```

```

if (signum == SIGINT || signum == SIGTERM) {
    printf("\nCaught signal %d, shutting down...\n", signum);
    force_quit = true;
}

}

int main(int argc, char **argv) {
    struct rte_mempool *mbuf_pool;
    int ret;

    // Initialize DPDK EAL
    ret = rte_eal_init(argc, argv);
    if (ret < 0)
        rte_exit(EXIT_FAILURE, "EAL init failed\n");

    signal(SIGINT, signal_handler);
    signal(SIGTERM, signal_handler);

    // Check ports (need at least 2)
    uint16_t nb_ports = rte_eth_dev_count_avail();
    if (nb_ports < 2)
        rte_exit(EXIT_FAILURE, "Need at least 2 ports\n");

    // Create mbuf pool
    mbuf_pool = rte_pktmbuf_pool_create("MBUF_POOL", NUM_MBUFS * nb_ports,
                                        MBUF_CACHE_SIZE, 0, RTE_MBUF_DEFAULT_BUF_SIZE,

```

```

        rte_socket_id());

if (!mbuf_pool)
    rte_exit(EXIT_FAILURE, "Cannot create mbuf pool\n");

// Configure & start ports (port 0 = RX, port 1 = TX)

struct rte_eth_conf port_conf = {0};

for (uint16_t port = 0; port < 2; port++) {
    ret = rte_eth_dev_configure(port, 1, 1, &port_conf);
    if (ret < 0)
        rte_exit(EXIT_FAILURE, "Cannot configure port %u\n", port);

    ret = rte_eth_rx_queue_setup(port, 0, RX_RING_SIZE,
                               rte_eth_dev_socket_id(port), NULL, mbuf_pool);
    if (ret < 0)
        rte_exit(EXIT_FAILURE, "rx queue setup failed on port %u\n", port);

    ret = rte_eth_tx_queue_setup(port, 0, TX_RING_SIZE,
                               rte_eth_dev_socket_id(port), NULL);
    if (ret < 0)
        rte_exit(EXIT_FAILURE, "tx queue setup failed on port %u\n", port);

    rte_eth_dev_start(port);
    rte_eth_promiscuous_enable(port);
}

// Initialize HFSC scheduler

```

```

hfsc_init();

printf("HFSC scheduler initialized.\n");
printf("Forwarding: RX port 0 → HFSC → TX port 1\n");

// Main poll loop

while (!force_quit) {

    // RX from port 0

    struct rte_mbuf *rx_pkts[BURST_SIZE];

    uint16_t nb_rx = rte_eth_rx_burst(0, 0, rx_pkts, BURST_SIZE);

    for (uint16_t i = 0; i < nb_rx; i++) {

        if (hfsc_packet_in(rx_pkts[i]) != 0) {

            rte_pktmbuf_free(rx_pkts[i]);

        }

    }

    // Dequeue from HFSC & TX to port 1

    struct rte_mbuf *tx_pkts[BURST_SIZE];

    uint16_t nb_tx = 0;

    while (nb_tx < BURST_SIZE) {

        struct rte_mbuf *m = hfsc_packet_out();

        if (!m) break;

        tx_pkts[nb_tx++] = m;

    }

}

```

```
if (nb_tx > 0) {
    uint16_t nb_sent = rte_eth_tx_burst(1, 0, tx_pkts, nb_tx);
    for (uint16_t i = nb_sent; i < nb_tx; i++) {
        rte_pktmbuf_free(tx_pkts[i]);
    }
}

// Cleanup
printf("Shutting down...\n");
for (uint16_t port = 0; port < 2; port++) {
    rte_eth_dev_stop(port);
}

rte_eal_cleanup();
return 0;
}

//*****
# SPDX-License-Identifier: BSD-3-Clause
# Copyright(c) 2026 Manikant

include $(RTE_SDK)/mk/rte.vars.mk

APP = dpdk-hfsc_forward
```

```
SRCS-y += main.c
```

```
SRCS-y += hfsc.c
```

```
CFLAGS += -O3 -g -I.
```

```
LIBS += -l rte_eal -l rte_mbuf -l rte_mempool_ring -l rte_ring -l rte_ethdev
```

```
LIBS += -l rte_bus_pci -l rte_pci -l rte_kvargs -l dl -l pthread -l numa
```

```
include $(RTE_SDK)/mk/rte.extapp.mk
```

```
*****
```

```
# Build (meson/ninja - modern DPDK style)
```

```
meson setup build
```

```
ninja -C build
```

```
# Or classic make (if your DPDK version uses it)
```

```
make -C examples/hfsc_forward T=x86_64-native-linux-gcc
```

```
# Run (very important: use --file-prefix to avoid conflicts)
```

```
sudo ./build/examples/dpdk-hfsc_forward \
```

```
-l 0-3 \
```

```
-n 4 \
```

```
--huge-dir=/mnt/huge \
```

```
--file-prefix=hfsc \
```

```
--proc-type=primary
```

```
ninja -C build # or make as above  
sudo ./build/examples/dpdk-hfsc_forward \  
-l 0-3 -n 4 --huge-dir=/mnt/huge --file-prefix=hfsc
```